SENIOR STAFF ENGINEER

# Maya Patel

maya.patel@example.com • +1 (415) 555-0142 • San Francisco, CA • mayapatel.ai • github.com/mpatel • linkedin.com/in/mayapatel

## SUMMARY

Senior staff engineer with 10+ years shipping ML infrastructure at frontier-lab scale. Architected training and inference systems serving 200M+ users. Led cross-org platform bets from design doc to GA, grew teams from 4 to 30, and drove $40M+ in annual compute savings. Two best-paper nominations (NeurIPS, ICML).

## EXPERIENCE

### Senior Staff Engineer
2022 – PRESENT

OpenAI · San Francisco

- Architected the next-generation training platform that reduced GPT-scale pre-training cost by 38%, saving $40M+ annually across 50k+ H100 GPU-hours per run.
- Led a cross-functional team of 18 engineers to ship a real-time inference router that serves 200M+ monthly active users at 99.97% uptime.
- Designed and rolled out the post-training evaluation framework adopted by 60+ research teams; caught 3 critical regressions before public launch.
- Promoted from Staff to Senior Staff within 18 months; nominated for the internal Distinguished Engineer track.

### Staff Research Engineer
2020 – 2022

Anthropic · San Francisco

- Built the distributed inference stack behind Claude from prototype to production; achieved 3.2x throughput improvement at 45% lower p99 latency.
- Open-sourced an interpretability toolkit that reached 12k GitHub stars in 8 months; adopted by 5 external research labs.
- Spearheaded the RLHF data pipeline serving 500k+ preference comparisons per day with sub-second freshness.

### Senior Software Engineer
2017 – 2020

Google DeepMind · Mountain View

- Shipped the unified training telemetry platform used by 400+ researchers across the Brain–DeepMind merger, reducing experiment debug time by 60%.
- First-authored a NeurIPS 2020 paper on efficient attention at 64k sequence length (best paper honorable mention, 280+ citations).
- Led the Gemini-precursor multi-modal data pipeline processing 2B+ image-text pairs per day.

### Software Engineer
2015 – 2017

Meta AI (FAIR) · Menlo Park

- Built the PyTorch distributed training backend that scaled model-parallel training to 256 GPUs with near-linear efficiency.
- Co-authored 3 peer-reviewed papers on large-scale representation learning (CVPR, ICLR).

## PROJECTS

### nanoeval
2024

Creator & maintainer · github.com/mpatel/nanoeval

- A minimal, typed eval framework for LLMs. Adopted by OpenAI, Anthropic, and Cohere for internal benchmarking. 8.2k ★ on GitHub.

### Triton Cookbook
2023

Open-source lead · github.com/mpatel/triton-cookbook

- Annotated GPU kernel recipes (attention, RMSNorm, MoE routing) used in 20+ university ML courses. 15k ★.

## EDUCATION

### Stanford University
2013 – 2015

M.S. Computer Science · Stanford, CA

Thesis: "Distributed Training of Sequence Models at Scale" · Advisor: Prof. Christopher Ré · Dean's List

### UC Berkeley
2009 – 2013

B.S. Electrical Engineering & Computer Science · Berkeley, CA

GPA 3.96 · Phi Beta Kappa · Regents' Scholar · ACM ICPC regional finalist

## SKILLS

**Languages:** Python, CUDA, C++, Rust, Triton, Go, TypeScript

**ML & systems:** PyTorch, JAX, Ray, Megatron-LM, vLLM, DeepSpeed, NCCL, TensorRT

**Infrastructure:** H100/A100 clusters, Slurm, Kubernetes, AWS, GCP, Kafka, PostgreSQL, Redis

**Leadership:** Technical roadmapping, architecture review, hiring (100+ interviews), mentoring, cross-org alignment